Lecture 20 (B19): Continuous-Time Heterogeneous Agents — Theory

Continuous-Time Heterogeneous Agent Models: Theory

Simon Scheidegger

HEC, University of Lausanne

Based on: Achdou, Han, Lasry, Lions & Moll (2022), *Rev. Econ. Stud.* 89(1), 45–86.

Gu, Lauriere, Merkel & Payne (2024), "Global Solutions to Master Equations." arXiv:2406.13726.

Payne, J. (2023–2024), ECO 529 Lecture Notes, Princeton University.

Materials: https://github.com/sischei/Deep_Learning_for_Solving_And_Estimating_Dynamic_Economic_Models

Where We Are in This Course

DL Basics → DEQNs → IRBC, NAS → OLG, Young's → Agentic toolkit → **Now** PINNs + CT HA → Surrogates → Climate

Block 1 (PINNs):

- Single PDEs: Poisson, HJB, Black–Scholes
- Neural network $\approx$ PDE solution
- Residual loss + automatic differentiation

Block 2 (this block):

- Coupled PDEs: HJB + KFE
- PDEs arise from economic equilibrium
- Continuum of heterogeneous agents

From PINNs to Heterogeneous Agent Models

Block 1: PINNs
Solve *one* PDE
$\mathcal{L}[u] = 0$ on Ω

$\xrightarrow{\text{extend}}$

Block 2: Het. Agents
Solve *coupled* system
HJB + KFE + Market Clearing

Recall from Deck 1: the cake-eating HJB you just solved with a PINN is the single-PDE special case (one agent, no income, no distribution). We now *add* idiosyncratic income, which forces a population density g into the picture and brings the KFE alongside the HJB.

The key insight: in continuous-time macro with heterogeneous agents, equilibrium is characterized by a system of PDEs:

1. **HJB equation** (backward): each agent optimizes $\Rightarrow$ value function V, policy c^*
2. **Kolmogorov forward equation** (forward): population density g evolves given c^*
3. **Market clearing**: aggregate supply = aggregate demand $\Rightarrow$ prices r, w

Why Continuous Time?

Advantages of continuous-time formulation:

- Analytical tractability: Ito calculus provides closed-form FOCs
- Clean separation: HJB (backward) $\leftrightarrow$ KFE (forward)
- Powerful numerical methods: finite differences, PINNs
- No expectations operator: replaced by differential operators
- Connection to mean field games (Lasry & Lions, 2007)

Key References

- Achdou, Han, Lasry, Lions & Moll (2022), *REStud*
- Kaplan, Moll & Violante (2018), *AER*
- Moll lecture notes: `benjaminmoll.com/lectures/`

Discrete $\rightarrow$ Continuous

Earlier discrete-time lectures (DEQNs, Young's method) used *discrete time*. Today: *continuous time* $\Rightarrow$ PDEs.

The Aiyagari–Bewley–Huggett Framework

- Benchmark models in "quantitative" macroeconomics:
  - **Bewley** (1980, 1983): precautionary saving with borrowing constraints
  - **Huggett** (1993, *JEDC*): endowment economy, risk-free rate determination
  - **Aiyagari** (1994, *QJE*): production economy, general equilibrium
  - **Krusell & Smith** (1998, *JPE*): adds aggregate TFP shocks
- Core mechanism: idiosyncratic labor income risk is uninsurable
  - Households want to borrow when income is low $\rightarrow$ borrowing limits prevent this
  - Instead, households accumulate assets as precautionary savings
- These are building blocks for HANK models (Kaplan, Moll & Violante, 2018)
  - Heterogeneous Agent New Keynesian: monetary/fiscal policy with inequality

Goals for This Lecture

1. **Stochastic calculus refresher**: Brownian motion, Ito's lemma, Poisson processes
2. **Derive the Kolmogorov forward equation** (KFE / Fokker–Planck):
   - First as a standalone concept from physics (diffusion of particles)
   - Then in economics: continuum of agents $\rightarrow$ wealth distribution
3. **The Hamilton–Jacobi–Bellman equation**: stochastic optimal control in CT
4. **Competitive equilibrium**: coupled HJB + KFE + market clearing
5. **The master equation**: collapsing the system into a single PDE

Deck 2 (after coffee): *numerical methods* — finite differences, PINNs, and deep learning (EMINNs) for solving these coupled PDEs.

Table of Contents

Stochastic Calculus Refresher

The Kolmogorov Forward Equation

The HJB Equation

Competitive Equilibrium

The Master Equation

Key References

Foundational models:

- Huggett, M. (1993). "The Risk-Free Rate in Heterogeneous-Agent Incomplete-Insurance Economies." *JEDC* 17(5–6), 953–969.
- Aiyagari, S.R. (1994). "Uninsured Idiosyncratic Risk and Aggregate Saving." *QJE* 109(3), 659–684.
- Krusell, P. & Smith, A.A. (1998). "Income and Wealth Heterogeneity in the Macroeconomy." *JPE* 106(5), 867–896.

Continuous-time formulations:

- Achdou, Han, Lasry, Lions & Moll (2022). "Income and Wealth Distribution in Macroeconomics: A Continuous-Time Approach." *REStud* 89(1), 45–86.
- Gu, Lauriere, Merkel & Payne (2024). "Global Solutions to Master Equations for CT Het. Agent Models." arXiv:2406.13726.
- Lasry & Lions (2007). "Mean Field Games." *Jpn. J. Math.* 2(1), 229–260.

Part I

Stochastic Calculus Refresher

Brownian Motion (Wiener Process)

Definition

A stochastic process $\{B_t\}_{t\geq 0}$ is a standard Brownian motion if:

1. $B_0 = 0$
2. **Independent increments**: $B_t - B_s \perp B_s - B_r$ for $r < s < t$
3. **Gaussian**: $B_t - B_s \sim \mathcal{N}(0, t-s)$
4. **Continuous paths**: $t \mapsto B_t$ is continuous a.s.

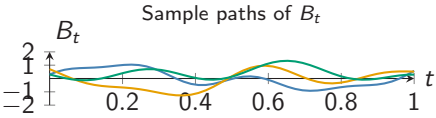

Key properties:

- $\mathbb{E}[B_t] = 0, \quad \mathrm{Var}(B_t) = t$
- Paths are nowhere differentiable
- Quadratic variation: $\langle B \rangle_t = t$

Brownian Motion: Random Walk Limit

Discrete-time intuition: Brownian motion is the limit of a scaled random walk.

Let $\Delta t > 0$ and define $X_t^{\Delta t}$ by steps $\pm\sqrt{\Delta t}$ with equal probability:

$$X_{t+\Delta t}^{\Delta t} = X_t^{\Delta t} + \sqrt{\Delta t}\,\varepsilon_t, \qquad \varepsilon_t \overset{\text{iid}}{\sim} \{-1, +1\} \text{ with prob. } \tfrac{1}{2} \text{ each.}$$

As $\Delta t \to 0$: $\quad X_t^{\Delta t} \xrightarrow{\text{d}} B_t \qquad$ (Donsker's theorem).

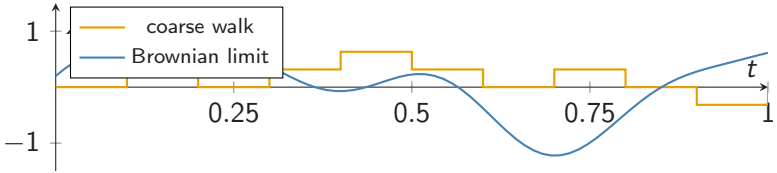

The $\sqrt{\Delta t}$ scaling is crucial: it ensures $\mathrm{Var}(X_t) = t$, not 0 or ∞.

Itô Processes and SDEs

General Itô Process

A stochastic process X_t follows an Itô process if:

$$\boxed{dX_t = \mu(X_t, t)\,\mathrm{d}t + \sigma(X_t, t)\,\mathrm{d}B_t}$$

- $\mu(X_t, t)$: drift (deterministic trend)
- $\sigma(X_t, t)$: diffusion coefficient (volatility / noise intensity)
- B_t: standard Brownian motion

Important examples:

Geometric Brownian Motion (GBM)

$dS_t = \mu S_t\,\mathrm{d}t + \sigma S_t\,\mathrm{d}B_t$

Used for: stock prices, GDP

Ornstein–Uhlenbeck (OU) Process

$dZ_t = \eta(\bar{Z} - Z_t)\,\mathrm{d}t + \sigma\,\mathrm{d}B_t$

Used for: productivity, interest rates

The Ornstein–Uhlenbeck Process

$$dZ_t = \eta(\bar{Z} - Z_t)\,\mathrm{d}t + \sigma\,\mathrm{d}B_t$$

Properties:

- Mean-reverting: Z_t is pulled toward $\bar{Z}$ at rate η
- Stationary distribution: $Z_\infty \sim \mathcal{N}\left(\bar{Z},\ \frac{\sigma^2}{2\eta}\right)$
- Closed-form solution by integrating factor.

In this lecture:

- Aggregate TFP: $dZ_t = \eta(\bar{Z} - Z_t)\,\mathrm{d}t + \sigma\,\mathrm{d}B_t^0$
- Krusell–Smith (1998): Z_t drives aggregate fluctuations

OU sample paths ($\bar{Z} = 1$)

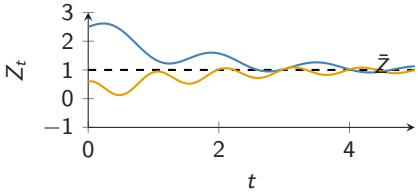

Poisson Process and Continuous-Time Markov Chains

Poisson Process

N_t: counting process with intensity $\lambda > 0$.

- Jumps of size $+1$ at random times
- Inter-arrival times $\sim \text{Exp}(\lambda)$
- $\Pr(\text{jump in } [t, t+h)) = \lambda h + o(h)$

Income switching (for our models):

- Labor state $n_t \in \{n_1, n_2\}$ (e.g., employed/unemployed)
- $n_1 \to n_2$ with intensity λ_1
- $n_2 \to n_1$ with intensity λ_2
- Ergodic probabilities: $\pi_1 = \frac{\lambda_2}{\lambda_1+\lambda_2}$, $\pi_2 = \frac{\lambda_1}{\lambda_1+\lambda_2}$

Income state path

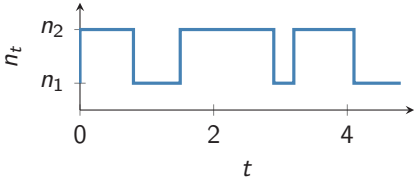

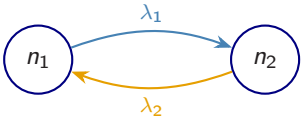

Itô's Lemma: The Chain Rule of Stochastic Calculus

Itô's Lemma

Let X_t follow $dX_t = \mu\,\mathrm{d}t + \sigma\,\mathrm{d}B_t$ and let $f \in C^2$. Then:

$$\boxed{df(X_t) = \left[f'(X_t)\,\mu + \tfrac{1}{2}f''(X_t)\,\sigma^2\right]\,\mathrm{d}t + f'(X_t)\,\sigma\,\mathrm{d}B_t}$$

The key difference from ordinary calculus:

- In ordinary calculus: $df = f'\,dX$ (first-order only)
- In stochastic calculus: $df = f'\,dX + \frac{1}{2}f''\,(dX)^2$ (second-order term survives!)
- Why? Because $(dB_t)^2 = dt \neq 0$ (Brownian motion has "infinite" variation)

Differential algebra rules:

$$\mathrm{d}t \cdot \mathrm{d}t = 0, \qquad \mathrm{d}t \cdot \mathrm{d}B_t = 0, \qquad \mathrm{d}B_t \cdot \mathrm{d}B_t = \mathrm{d}t$$

Itô's Lemma: Worked Example

Example: Let $dX_t = \mu X_t \, \mathrm{d}t + \sigma X_t \, \mathrm{d}B_t$ (GBM). Find $d(\ln X_t)$.

Step 1: Let $f(x) = \ln x$, so $f'(x) = 1/x$, $f''(x) = -1/x^2$.

Step 2: Apply Itô's Lemma with $\mu(X_t) = \mu X_t$, $\sigma(X_t) = \sigma X_t$:

$$d(\ln X_t) = \frac{1}{X_t} \cdot \mu X_t \, \mathrm{d}t + \frac{1}{2}\left(-\frac{1}{X_t^2}\right)(\sigma X_t)^2 \, \mathrm{d}t + \frac{1}{X_t} \cdot \sigma X_t \, \mathrm{d}B_t$$

$$= \boxed{\left(\mu - \tfrac{1}{2}\sigma^2\right) \, \mathrm{d}t + \sigma \, \mathrm{d}B_t}$$

Insight: $\ln X_t$ is an Itô process with *constant* drift $\mu - \frac{1}{2}\sigma^2$ and diffusion σ.

$\Rightarrow$ Solution: $X_t = X_0 \exp\left[\left(\mu - \frac{1}{2}\sigma^2\right) t + \sigma B_t\right]$

The $-\frac{1}{2}\sigma^2$ Itô correction is why expected return $\neq$ log return for volatile assets.

Itô's Lemma: Multivariate Version

For $V(X_t, t)$ where $dX_t = \mu\,\mathrm{d}t + \sigma\,\mathrm{d}B_t$:

$$\boxed{dV = \left[\frac{\partial V}{\partial t} + \mu\frac{\partial V}{\partial x} + \frac{1}{2}\sigma^2\frac{\partial^2 V}{\partial x^2}\right]\mathrm{d}t + \sigma\frac{\partial V}{\partial x}\,\mathrm{d}B_t}$$

With jumps (Poisson process with intensity λ, jump ΔX):

$$dV = \underbrace{\left[\frac{\partial V}{\partial t} + \mu\frac{\partial V}{\partial x} + \frac{1}{2}\sigma^2\frac{\partial^2 V}{\partial x^2}\right]\mathrm{d}t + \sigma\frac{\partial V}{\partial x}\,\mathrm{d}B_t}_{\text{diffusion part}} + \underbrace{\left[V(X_{t^-} + \Delta X, t) - V(X_{t^-}, t)\right]\mathrm{d}N_t}_{\text{jump part}}$$

We will use the multivariate version repeatedly:
HJB derivation (Part IV) and **KFE derivation** (Part II).

Part II

The Kolmogorov Forward Equation

(Fokker–Planck Equation)

What Is the Kolmogorov Forward Equation?

The Kolmogorov forward equation (KFE), also known as the Fokker–Planck equation, describes how the **probability density** of a stochastic process evolves over time.

Origin: physics and applied mathematics.

- **A.N. Kolmogorov** (1931): general theory of continuous Markov processes
- **A.D. Fokker** (1914) & **M. Planck** (1917): diffusion of particles
- Applications: heat conduction, fluid dynamics, plasma physics, chemical kinetics, population genetics, neuroscience, ...

A.N. Kolmogorov, Wikimedia Commons

Key idea: micro dynamics (drift + noise) $\Rightarrow$ macro density law of motion.

Standalone Example: Diffusion of Particles

Setup (no economics!): Consider a large population of particles, each independently following:

$$dX_t = \mu \, \mathrm{d}t + \sigma \, \mathrm{d}B_t$$

where μ is a constant drift and σ is a constant diffusion coefficient.

Question: If at time $t = 0$ the particles have density $g_0(x)$, what is the density $g(x, t)$ at time $t > 0$?

Kolmogorov Forward Equation (constant coefficients)

$$\boxed{\frac{\partial g}{\partial t}(x, t) = -\mu \frac{\partial g}{\partial x}(x, t) + \frac{\sigma^2}{2} \frac{\partial^2 g}{\partial x^2}(x, t)}$$

Two competing effects:

- $-\mu \, \partial_x g$: advection — drift transports the density
- $\frac{\sigma^2}{2} \partial_{xx} g$: diffusion — noise spreads the density

Deriving the KFE from First Principles

Goal: Derive $\partial_t g = -\mu \partial_x g + \frac{\sigma^2}{2} \partial_{xx} g$ from basic probability.
Step 1 — For any smooth test function $\varphi(x)$: $\mathbb{E}[\varphi(X_t)] = \int_{-\infty}^{\infty} \varphi(x)\, g(x,t)\, \mathrm{d}x$
Step 2 — Differentiate w.r.t. t and use Itô's lemma on $\varphi(X_t)$:

$$\frac{d}{dt} \int \varphi\, g\, \mathrm{d}x = \mathbb{E}\Big[\mu\, \varphi'(X_t) + \tfrac{\sigma^2}{2} \varphi''(X_t)\Big] = \int \Big[\mu\, \varphi' + \tfrac{\sigma^2}{2} \varphi''\Big]\, g\, \mathrm{d}x$$

Step 3 — Integrate by parts (assuming $g \to 0$ at $\pm\infty$):

$$\int \mu\, \varphi'\, g\, \mathrm{d}x = -\int \varphi\, \mu\, \partial_x g\, \mathrm{d}x \qquad \text{(once)}$$

$$\int \tfrac{\sigma^2}{2} \varphi''\, g\, \mathrm{d}x = \int \varphi\, \tfrac{\sigma^2}{2} \partial_{xx} g\, \mathrm{d}x \qquad \text{(twice)}$$

Step 4 — Since this holds for *all* test functions φ: $\boxed{\partial_t g = -\mu\, \partial_x g + \frac{\sigma^2}{2}\, \partial_{xx} g}$ ✓

The KFE: General Itô Diffusion

For $dX_t = \mu(X_t)\,\mathrm{d}t + \sigma(X_t)\,\mathrm{d}B_t$ with state-dependent drift and diffusion, the density $g(x,t)$ satisfies the **general KFE**:

Kolmogorov Forward / Fokker–Planck Equation

$$\boxed{\frac{\partial g}{\partial t}(x,t) = -\frac{\partial}{\partial x}\left[\mu(x)\,g(x,t)\right] + \frac{1}{2}\frac{\partial^2}{\partial x^2}\left[\sigma(x)^2\,g(x,t)\right]}$$

Important: when σ depends on x, it enters *inside* the second derivative:
$\frac{1}{2}\partial_{xx}\left[\sigma(x)^2 g\right] \neq \frac{\sigma(x)^2}{2}\partial_{xx} g$

In divergence form: $\partial_t g = -\partial_x \underbrace{\left[\mu(x)g - \frac{1}{2}\partial_x(\sigma^2 g)\right]}_{=:\text{probability flux } J(x,t)}$ Conservation law: $\partial_t g + \partial_x J = 0$.

Example 1: The Heat Equation (Pure Diffusion)

Setting: No drift ($\mu = 0$), constant diffusion σ. Particles start at x_0: $\partial_t g = \frac{\sigma^2}{2}\partial_{xx} g$, $g(x, 0) = \delta(x - x_0)$.

Solution: Gaussian with spreading variance: $g(x, t) = \frac{1}{\sqrt{2\pi\sigma^2 t}} \exp\left(-\frac{(x-x_0)^2}{2\sigma^2 t}\right)$

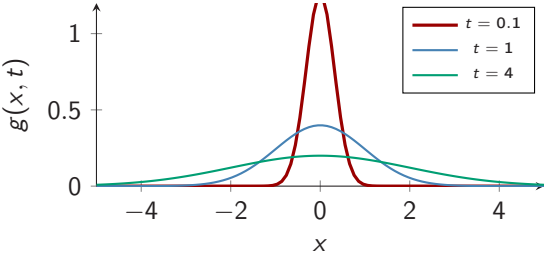

Physical interpretation:

- Heat spreads from a point source
- Variance grows linearly: $\text{Var} = \sigma^2 t$
- Total "mass" conserved: $\int g = 1$
- Most fundamental PDE in physics

This is exactly what PINNs solve!

Example 2: Ornstein–Uhlenbeck Process

Setting: $dX_t = \theta(\bar{x} - X_t)\,\mathrm{d}t + \sigma\,\mathrm{d}B_t$ (mean-reverting).

KFE: $\partial_t g = \theta\,\partial_x[(x - \bar{x})\,g] + \frac{\sigma^2}{2}\partial_{xx} g$

Stationary distribution ($\partial_t g = 0$):

$$\boxed{g^*(x) = \frac{1}{\sqrt{2\pi\sigma^2/(2\theta)}} \exp\left(-\frac{(x-\bar{x})^2}{2\cdot\sigma^2/(2\theta)}\right) = \mathcal{N}\left(\bar{x},\ \frac{\sigma^2}{2\theta}\right)}$$

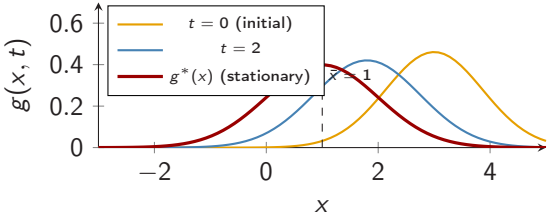

Interpretation:

- Mean reversion (θ) concentrates mass around $\bar{x}$
- Diffusion (σ) spreads mass
- Balance $\Rightarrow$ Gaussian steady state
- Larger $\theta \Rightarrow$ tighter distribution

This will model aggregate productivity Z_t in the Krusell–Smith model.

Example 3: Boundaries and Borrowing Constraints

Setting: $dX_t = \mu\,\mathrm{d}t + \sigma\,\mathrm{d}B_t$ on $X_t \geq 0$.

- **Absorbing BC**: mass exits at the boundary.
- **Reflecting BC**: no net probability flow across the boundary.
- In Huggett/Aiyagari (asset drift only), borrowing constraint $a \geq \underline{a}$ is a state constraint.

For the absolutely continuous density part:

$$J(a, n) := s(a, n)\, g_t(a, n), \qquad J(\underline{a}, n) = 0.$$

Boundary atoms can form at $\underline{a}$ when constrained agents pile up.

No-outflow at boundary

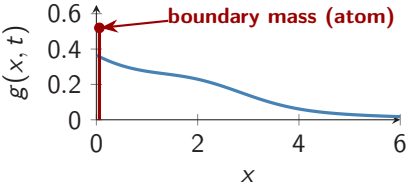

From Physics to Economics: A Continuum of Agents

Key insight (propagation of chaos):
Consider a continuum of agents $i \in [0, 1]$, each independently following:

$$dX_t^i = \mu(X_t^i)\,\mathrm{d}t + \sigma(X_t^i)\,\mathrm{d}B_t^i$$

where B_t^i are *independent* Brownian motions (idiosyncratic shocks).
By the **law of large numbers** at the population level:

- The cross-sectional density $g(x, t)$ evolves deterministically
- Even though each individual path is random!
- The density satisfies the KFE: $\partial_t g = -\partial_x[\mu g] + \frac{1}{2}\partial_{xx}[\sigma^2 g]$

Analogy: Each molecule of air moves randomly (Brownian motion), but the *temperature distribution* in a room evolves deterministically (heat equation).

In economics: each household's income is random, but the *wealth distribution* evolves deterministically (KFE).

KFE in Economics: The Wealth Distribution

Setup: Agent i has wealth a_t^i and income state $n_t^i \in \{n_1, n_2\}$.
Individual wealth evolves as:

$$da_t^i = \underbrace{s(a_t^i, n_t^i)}_{\text{savings} = \text{income} - \text{consumption}} \, \mathrm{d}t, \qquad a_t^i \geq \underline{a}$$

Income n_t^i switches between $n_1 \leftrightarrow n_2$ with Poisson intensities λ_1, λ_2.
The cross-sectional density $g_t(a, n)$ satisfies:

$$\boxed{\partial_t g_t(a, n) = -\partial_a[s(a, n)\, g_t(a, n)] - \lambda(n)\, g_t(a, n) + \lambda(\hat{n})\, g_t(a, \hat{n})}$$

where $\hat{n}$ denotes the complementary income state.

- $-\partial_a[s \cdot g]$: flow of agents along the wealth axis (savings)
- $-\lambda(n)\, g(a, n)$: agents leaving state n (exit from $n \to \hat{n}$)
- $+\lambda(\hat{n})\, g(a, \hat{n})$: agents entering state n (entry from $\hat{n} \to n$)

Stationary Wealth Distribution

Setting $\partial_t g = 0$, the stationary (ergodic) distribution solves:

$$
\begin{aligned}
0 = & -\partial_a[s(a,n)\, g^*(a,n)] \\
& -\lambda(n)\, g^*(a,n) + \lambda(\hat{n})\, g^*(a,\hat{n})
\end{aligned}
$$

Key features of $g^*(a,n)$:

- Right-skewed: many poor, few rich
- Mass point at borrowing constraint $\underline{a}$ for low-income agents
- Separate densities for each income state
- Normalization: $\sum_n \int_{\underline{a}}^{\infty} g^*(a,n)\, \mathrm{d}a = 1$

This is the object solved numerically in Aiyagari (1994) and Achdou et al. (2022).

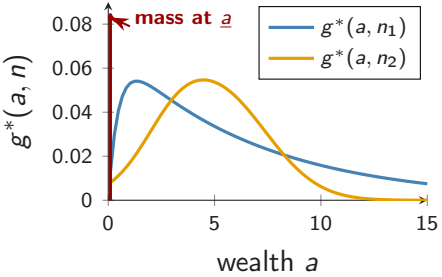

Why Does the KFE Not Have a dB_t^0 Term?

Observation: The KFE for $g_t(a, n)$ is a deterministic PDE, even though each agent faces random shocks.

Why?

- Agents face idiosyncratic shocks dB_t^i (independent across i)
- By the **law of large numbers**: randomness washes out at the population level
- $\Rightarrow$ The density g_t evolves deterministically

When would dB_t^0 appear in the KFE?

- If there are aggregate shocks (e.g., TFP Z_t following an OU process)
- Then prices r_t, w_t become stochastic $\Rightarrow$ drift is stochastic
- $\Rightarrow$ The density g_t becomes a stochastic process adapted to $\mathcal{F}_t^0$
- This is the Krusell–Smith (1998) setting — much harder!

Aiyagari (no aggregate shocks): KFE is deterministic ✓ **Krusell–Smith** (aggregate shocks): KFE is stochastic $\Rightarrow$ "master equation"

Connection to Mean Field Games

The KFE is one half of the Mean Field Game (MFG) system:

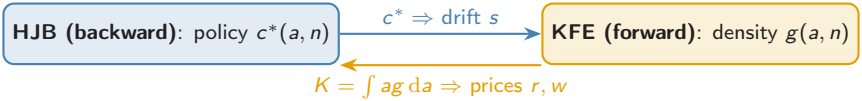

- **Lasry & Lions** (2007): formalized MFGs as coupled PDE systems
- In economics: "Mean Field Game" = competitive equilibrium with continuum of agents
- Each agent is "small" $\Rightarrow$ takes distribution as given
- But collectively, agents *determine* the distribution
- Fixed point: beliefs about g are consistent with optimal behavior

KFE: Summary

Version	KFE	Application
Pure diffusion	$\partial_t g = -\mu \partial_x g + \frac{\sigma^2}{2} \partial_{xx} g$	Heat equation, particle diffusion
General Itô	$\partial_t g = -\partial_x [\mu g] + \frac{1}{2} \partial_{xx} [\sigma^2 g]$	OU process, GBM
With Poisson	$\cdots - \lambda(n) g(n) + \lambda(\hat{n}) g(\hat{n})$	Income switching
Economic	drift $= s(a, n)$ (endogenous)	Wealth distribution

Key takeaways:

- The KFE maps $\{\text{drift } \mu, \text{diffusion } \sigma, \text{initial density } g_0\} \to$ density evolution $g(x, t)$
- In economics: drift = savings (endogenous, from HJB)
- Without aggregate shocks: KFE is deterministic (LLN)
- With aggregate shocks: KFE becomes stochastic $\Rightarrow$ "master equation"
- State constraints can require boundary atoms in addition to an interior density

Part III

The Hamilton–Jacobi–Bellman Equation

Stochastic Optimal Control in Continuous Time

The Agent's Optimization Problem

Setup: Agent i with state $x_t^i = (a_t^i, n_t^i)$ chooses consumption c_t^i to solve:

$$\max_{\{c_t^i\}_{t\geq 0}} \mathbb{E}\left[\int_0^\infty e^{-\rho t}\, u(c_t^i)\,\mathrm{d}t\right]$$

subject to:

$$da_t^i = \underbrace{(\tilde{w}_t n_t^i + \tilde{r}_t a_t^i - c_t^i)}_{\equiv\, s(c_t^i,\, a_t^i,\, n_t^i,\, \tilde{w}_t,\, \tilde{r}_t)}\,\mathrm{d}t, \qquad a_t^i \geq \underline{a}$$

$$n_t^i \in \{n_1, n_2\}, \quad \text{switches with Poisson intensity } \lambda(n_t^i)$$

- $u(c) = \frac{c^{1-\gamma}}{1-\gamma}$: CRRA utility, $\gamma > 0$ (risk aversion)
- $\rho > 0$: discount rate
- $(\tilde{r}, \tilde{w})$: agent's beliefs about price processes (rental rate, wage)
- $\underline{a}$: exogenous borrowing limit

The Value Function

Definition

The value function $V(a, n)$ is the maximum expected discounted utility starting from state (a, n):

$$V(a, n) = \max_{\{c_t\}} \mathbb{E}\left[\int_0^\infty e^{-\rho t}\, u(c_t)\, \mathrm{d}t \;\middle|\; a_0 = a,\ n_0 = n\right]$$

Properties of V:

- V is increasing in a (more wealth $\Rightarrow$ higher utility)
- V is concave in a (diminishing marginal value of wealth)
- $V(a, n_2) > V(a, n_1)$ if $n_2 > n_1$ (higher income $\Rightarrow$ higher value)
- $V'(a) = \partial_a V(a, n)$ is the marginal value of wealth

Goal: Find $V(a, n)$ and the optimal policy $c^*(a, n)$. $\Rightarrow$ Use the Hamilton–Jacobi–Bellman equation.

Deriving the HJB: Five Steps

Step 1 — Dynamic Programming Principle. For small $h > 0$:

$$V(a,n) = \max_c \left\{ \int_0^h e^{-\rho t} u(c_t)\,\mathrm{d}t + e^{-\rho h}\,\mathbb{E}[V(a_h, n_h) \mid a_0 = a, n_0 = n] \right\}$$

Step 2 — Itô's lemma on $V(a_t, n_t)$ (between jumps of n):

$$dV = V'(a,n) \cdot s(c,a,n)\,\mathrm{d}t$$

(no diffusion term since $da = s\,\mathrm{d}t$ is deterministic between jumps)
Step 3 — Take expectations, account for Poisson jumps:

$$\mathbb{E}[dV] = \left[V'(a,n) \cdot s + \lambda(n)\big(V(a,\hat{n}) - V(a,n)\big)\right]\,\mathrm{d}t$$

Step 4 — Substitute into DPP, divide by h, let $h \to 0$:
Step 5 — Impose optimality over c:

$$\boxed{\rho V(a,n) = \max_c \left\{ u(c) + V'(a,n) \cdot (wn + ra - c) + \lambda(n)\big(V(a,\hat{n}) - V(a,n)\big) \right\}}$$

The HJB Equation: Interpretation

$$\rho V(a,n) = \max_c \Big\{ u(c) + \underbrace{V'(a,n)\cdot s(a,n,c)}_{\text{marginal value} \times \text{savings}} + \underbrace{\lambda(n)\big(V(a,\hat{n}) - V(a,n)\big)}_{\text{expected change from income switch}} \Big\}$$

Interpretation ("required return on the value function"):

- LHS: ρV = "required return" = discount rate $\times$ asset value
- RHS: "dividend" + "capital gain" = total return
  - $u(c)$: flow utility ("dividend")
  - $V'(a,n)\cdot s$: expected capital gain from saving
  - $\lambda(n)(V(\hat{n}) - V(n))$: expected gain/loss from income switch

The HJB is an asset pricing equation:
the value function is priced like a financial asset.

PAUSE — 5 min Finger Exercise: The HJB Equation

Exercise (5 min)

A household with wealth a earns income y and chooses consumption c to maximize utility. The HJB equation is:

$$\rho\, V(a) = \max_c \left\{ u(c) + V'(a)\,(y + ra - c) \right\}.$$

1. Take the first-order condition w.r.t. c. What is the relationship between $u'(c)$ and $V'(a)$?
2. With $u(c) = \log(c)$: express optimal c^* in terms of $V'(a)$.
3. Interpret: what does $V'(a)$ represent economically?

Solution: The HJB Equation

Answer

1. FOC: $u'(c^*) = V'(a)$. Marginal utility of consumption equals the marginal value of wealth.
2. With $u(c) = \log(c)$: $1/c^* = V'(a)$, so $c^*(a) = 1/V'(a)$.
3. $V'(a)$ is the shadow price of wealth—the marginal value of an extra unit of assets. When $V'(a)$ is high (wealth is very valuable), consumption is low (the agent saves aggressively).

Optimal Policy: The Consumption Function

First-order condition (differentiate HJB w.r.t. c):

$$u'(c^*) = V'(a,n) \qquad \Longleftrightarrow \qquad \boxed{c^*(a,n) = (u')^{-1}(V'(a,n))}$$

For CRRA utility $u(c) = \frac{c^{1-\gamma}}{1-\gamma}$ with $u'(c) = c^{-\gamma}$:

$$c^*(a,n) = [V'(a,n)]^{-1/\gamma}$$

Savings function:

$$s(a,n) = wn + ra - c^*(a,n)$$

State constraint at $a = \underline{a}$:

$$\begin{aligned} &0 < c^*(\underline{a},n) \leq wn + r\underline{a}, \\ &c^*(\underline{a},n) = \min\{[V_a(\underline{a},n)]^{-1/\gamma},\ wn + r\underline{a}\}, \\ &s(\underline{a},n) = wn + r\underline{a} - c^*(\underline{a},n) \geq 0. \end{aligned}$$

Interior: $u'(c^*) = V_a$. Boundary: constrain consumption so wealth cannot drift below $\underline{a}$.

Consumption and Savings: Qualitative Behavior

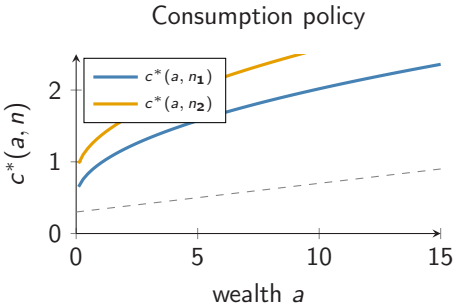

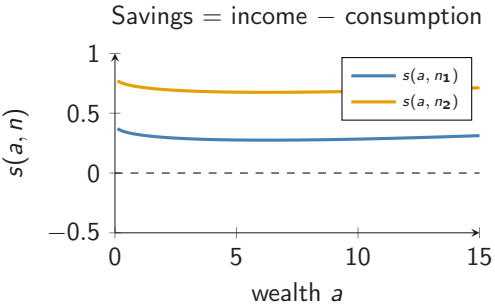

- Low-income agents near $\underline{a}$: savings near zero, often constrained
- High-income agents: positive savings $\Rightarrow$ accumulate wealth
- Precautionary savings: more saving than in the certainty benchmark

HJB with Aggregate Shocks (Krusell–Smith)

When aggregate TFP Z_t follows $dZ_t = \eta(\bar{Z} - Z_t)\,\mathrm{d}t + \sigma^z\,\mathrm{d}B_t^0$, the value function depends on the entire wealth distribution g: $V = V(a, n, z, g)$. The **HJB** becomes:

$$\begin{aligned}\rho V(a,n,z,g) = \max_c \Big\{ & u(c) + \partial_a V \cdot \big(w(z,g)\,n + r(z,g)\,a - c\big) \\ & + \lambda(n)\big(V(a,\hat{n},z,g) - V(a,n,z,g)\big) \\ & + \partial_z V \cdot \mu^z(z) + \tfrac{1}{2}(\sigma^z)^2\,\partial_{zz} V + \sum_j \int_{\underline{a}}^{\infty} \frac{\delta V}{\delta g_j}(a, n_j, z, g)(y) \cdot \mu_j^g(y,z,g)\,\mathrm{d}y \Big\}\end{aligned}$$

The Curse of Dimensionality

g is infinite-dimensional. Read $\frac{\delta V}{\delta g_j}(\cdot)(y)$ as the Gâteaux/functional derivative: how V changes when a small bump of mass is added at wealth y in productivity-state j. After a finite-dimensional approximation of g (cf. Deck 2 *Three Approximations*), it collapses to ordinary partials $\partial V / \partial \hat{\varphi}_i$. This makes the PDE intractable by standard methods $\Rightarrow$ deep learning (EMINNs — Deck 2).

The Coupled HJB–KFE System

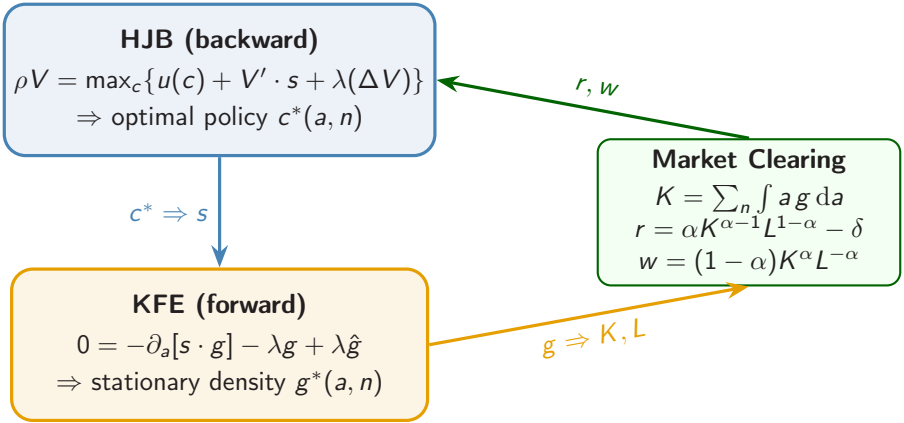

Solution method: iterate until convergence (Aiyagari, no aggregate shocks):

1. Guess $r \to$ solve HJB for $V, c^* \to$ solve KFE for $g^* \to$ compute $K \to$ update r

Part IV

Competitive Equilibrium

Aiyagari (1994) and Krusell–Smith (1998)

Setting: Agents and Preferences

- **Continuous time**, infinite horizon, one consumption good
- **Population**: continuum of agents $i \in I = [0, 1]$
- **Preferences**: $\mathbb{E}\left[\int_0^\infty e^{-\rho t}\, u(c_t^i)\, \mathrm{d}t\right], \quad u(c) = \frac{c^{1-\gamma}}{1-\gamma}$
- **Labor endowment**: $n_t^i \in \{n_1, n_2\}$ follows a CTMC:
  - $n_1 \to n_2$ with Poisson intensity λ_1; $\quad n_2 \to n_1$ with intensity λ_2

Parameter	**Interpretation**
ρ	Discount rate
γ	Risk aversion (inverse EIS)
$n_1 < n_2$	Low / high labor productivity
λ_1, λ_2	Transition intensities

Production Technology

Representative firm: $Y_t = e^{Z_t} F(K_t, L_t) = e^{Z_t} K_t^{\alpha} L_t^{1-\alpha}$

- K_t, L_t: capital and labor hired at time t
- F: CRS production function (Cobb–Douglas)
- **Aiyagari** (1994): $Z_t = \bar{Z}$ constant $\Rightarrow$ no aggregate risk
- **Krusell–Smith** (1998): Z_t follows OU: $dZ_t = \eta(\bar{Z} - Z_t)\,\mathrm{d}t + \sigma^z\,\mathrm{d}B_t^0$

Additional assumptions:

- Consumption $\leftrightarrow$ capital conversion at $1:1$ ratio
- Capital depreciates at rate δ

Markets

- **Incomplete financial markets**:
  - Agents cannot insure idiosyncratic income risk
  - Instead, agents "save and borrow claims on capital"
  - a_t: agent's claim on capital stock
  - Borrowing constraint: $a_t \geq \underline{a} > -\infty$
- **Competitive capital market** with rental rate r_t:

$$\text{Clearing:} \quad \int_I a_t^i \, \mathrm{d}i = K_t$$

- **Competitive labor market** with wage w_t:

$$\text{Clearing:} \quad \int_I n_t^i \, \mathrm{d}i = L_t$$

- **Goods market**: clears by Walras's Law

If $a_t < 0$, interpret this as borrowing from a frictionless intermediary that pools deposits and finances capital.

Firm Problem

The representative firm solves:

$$\max_{K_t, L_t} \left\{ e^{Z_t} F(K_t, L_t) - (r_t + \delta) K_t - w_t L_t \right\}$$

First-order conditions:

$$r_t = e^{Z_t} \partial_K F(K_t, L_t) - \delta = \alpha e^{Z_t} K_t^{\alpha - 1} L_t^{1-\alpha} - \delta$$

$$w_t = e^{Z_t} \partial_L F(K_t, L_t) = (1 - \alpha) e^{Z_t} K_t^{\alpha} L_t^{-\alpha}$$

Prices are functions of the aggregate state:

$$q_t = \begin{bmatrix} r_t \\ w_t \end{bmatrix} = Q(Z_t, g_t)$$

where $K_t = \sum_{j \in \{1,2\}} \int a \, g_t(a, n_j) \, \mathrm{d}a$.

Having explicit expressions for prices in terms of (Z_t, g_t) makes the problem tractable.

Kolmogorov Forward Equation in Equilibrium

Given stochastic processes $\{c^*, \tilde{r}, \tilde{w}\}$, at time $t > 0$, on $(\underline{a}, \infty)$, the density $g_t(a, n)$ satisfies the KFE:

$$\boxed{\partial_t g_t(a, n) = -\partial_a\Big[\big(\tilde{w}_t n + \tilde{r}_t a - c_t^*\big)\, g_t(a, n)\Big] - \lambda(n)\, g_t(a, n) + \lambda(\hat{n})\, g_t(a, \hat{n})}$$

Observations:

- Why does the KFE not have dB_t^0?
  - Idiosyncratic shocks wash out at the population level (LLN)
  - But with aggregate shocks (KS): g_t is stochastic via prices
- $g_t(a, n)$ is still stochastic in equilibrium since r and w depend on $\mathcal{F}_t^0$
- Low-income agents "stuck" at borrowing constraint $\Rightarrow$ boundary atom in g_t
- Technical: need measure-valued solution (Achdou et al., 2022)

Definition: Rational Expectations Equilibrium

Definition (Rational Expectations Equilibrium)

Given initial density g_0, an equilibrium is $\mathcal{F}_t^0$-adapted processes $\{c_t^i, g_t, z_t, q_t := [r_t, w_t]\}$ such that:

1. **Household optimality**: Given price belief $\tilde{q}$, c_t^i solves HJB.
2. **Firm optimality**: $r_t = e^{Z_t}\partial_K F - \delta, \quad w_t = e^{Z_t}\partial_L F$
3. **Market clearing**: $K_t = \sum_j \int a\, g_t(a, n_j)\, \mathrm{d}a, \quad L = \sum_j \int n_j\, g_t(a, n_j)\, \mathrm{d}a$
4. **Belief consistency**: $\tilde{q} = q \quad$ (beliefs = actual prices)

Derivation Map: Model $\Rightarrow$ HJB $\Rightarrow$ KFE $\Rightarrow$ Equilibrium

Step 1 (household control problem):

$$da_t = (wn_t + ra_t - c_t)\,\mathrm{d}t, \qquad n_t \in \{n_1, n_2\} \text{ (Poisson switching).}$$

Step 2 (HJB + FOC):

$$\rho V(a,n) = \max_c \left\{ u(c) + V_a(a,n)\,[wn + ra - c] + \lambda(n)\big(V(a,\hat{n}) - V(a,n)\big) \right\},$$

$$u'(c^*(a,n)) = V_a(a,n), \qquad s(a,n) := wn + ra - c^*(a,n).$$

Step 3 (KFE induced by optimal drift s):

$$0 = -\partial_a[s(a,n)g^*(a,n)] - \lambda(n)g^*(a,n) + \lambda(\hat{n})g^*(a,\hat{n}).$$

Step 4 (general equilibrium closure, Aiyagari):

$$K = \sum_n \int a\, g^*(a,n)\,\mathrm{d}a, \qquad L = \sum_n \int n\, g^*(a,n)\,\mathrm{d}a,$$

$$r = \alpha K^{\alpha-1} L^{1-\alpha} - \delta, \qquad w = (1-\alpha) K^{\alpha} L^{-\alpha}.$$

Fixed point: iterate on (r, w) until household supply and firm demand are consistent.

Three Notions of Equilibrium

Following Sargent (1999, 2008):

1. **Rational expectations equilibrium**
   - Agents' perceived law of motion = true law of motion
   - Even for data that would *never* be generated on the equilibrium path
2. **Self-confirming belief equilibrium**
   - Agents' model is the best statistical fit to data generated by the economy
   - Agents would never reject their model based on observed data
   - But may have incorrect off-equilibrium beliefs
3. **Optimal misspecified belief equilibrium**
   - Agents use an optimally chosen (but possibly wrong) model
   - Robust decision-making under model uncertainty

In this course: we focus on rational expectations. But note: neural network "solvers" can naturally accommodate bounded rationality (approximate equilibria).

Huggett (1993): Economic Mechanism

Environment

- Endowment economy with idiosyncratic income $y_t \in \{y_1, y_2\}$.
- Incomplete markets: only one risk-free bond b_t.
- Borrowing limit $b_t \geq \underline{b}$.
- Bond is in zero net supply.

Household trade-off

- Low income state: wants to borrow to smooth consumption.
- High income state: saves for self-insurance.
- Constraint creates a mass of constrained agents at $\underline{b}$.

How equilibrium price is determined

- For candidate return q, solve HJB $\Rightarrow$ policy $s(b, y)$.
- Propagate policy through KFE $\Rightarrow$ stationary density $g^*(b, y)$.
- Compute aggregate bond demand:

$$B^s(q) = \sum_y \int_{\underline{b}}^{\infty} b\, g^*(b, y)\, \mathrm{d}b.$$

- Equilibrium return: q^* such that $B^s(q^*) = 0$.

Implication: the risk-free rate is pinned down by precautionary self-insurance demand, not by a production FOC.

Aiyagari (1994): Economic Mechanism I

What changes vs. Huggett

- Add a representative firm with Cobb–Douglas production.
- Household asset is now a claim on aggregate capital a_t.
- Same incomplete insurance and borrowing constraint logic.
- Aggregate capital is endogenously positive: $K > 0$.

Supply side (households)

- For given (r, w), HJB gives policy $s(a, z)$.
- KFE gives density $g^*(a, z)$.
- Capital supplied by households: $K^s(r) = \sum_z \int_{\underline{a}}^{\infty} a\, g^*(a, z)\, \mathrm{d}a$.

Intuition

- Low productivity states generate demand for insurance via saving.
- High productivity states accumulate assets.
- With many uninsured households, aggregate saving pressure is high.
- This pushes the equilibrium return on capital downward.

Aiyagari (1994): Economic Mechanism II

General-equilibrium closure

Firm demand for capital

$$r = \alpha K^{\alpha-1} L^{1-\alpha} - \delta$$

$$w = (1-\alpha) K^{\alpha} L^{-\alpha}$$

These conditions imply a decreasing capital-demand schedule $K^d(r)$.

Household capital supply

$$K^s(r) = \sum_z \int_{\underline{a}}^{\infty} a\, g^*(a,z)\, \mathrm{d}a,$$

where g^* is the stationary density induced by optimal savings.

Equilibrium condition

$$K^s(r^*) = K^d(r^*), \qquad w^* = w(r^*).$$

Economic interpretation

Aiyagari is a full GE model: prices are jointly determined by

- household self-insurance behavior (asset supply), and
- firm marginal productivity (asset demand).

Huggett (1993): Formal Stationary PDE System

Primitives: endowment economy, $b_t \geq \underline{b}$, $y_t \in \{y_1, y_2\}$ with Poisson switching. Bonds pay instantaneous return r (price-vs-return convention: r, not a bond price q).

$$db_t = (r\, b_t + y_t - c_t)\,\mathrm{d}t, \qquad \max_{\{c_t\}} \mathbb{E}[]\left[\int_0^\infty e^{-\rho t} u(c_t)\,\mathrm{d}t\right].$$

HJB (state y):

$$\rho V = \max_c \{u(c) + V_b(rb + y - c) + \lambda(y)(V(b, \hat{y}) - V(b, y))\}, \quad u'(c^*) = V_b, \quad s = rb + y - c^*.$$

KFE on $(\underline{b}, \infty)$:

$$0 = -\partial_b[s(b, y)g^*(b, y)] - \lambda(y)g^*(b, y) + \lambda(\hat{y})g^*(b, \hat{y}),$$

$$J(\underline{b}, y) = s(\underline{b}, y)g^*(\underline{b}, y) = 0, \qquad \sum_y \int_{\underline{b}}^\infty g^*(b, y)\,\mathrm{d}b = 1.$$

Bond market clearing:

$$\sum_y \int_{\underline{b}}^\infty b\, g^*(b, y)\,\mathrm{d}b = 0.$$

Unknowns: (V, g^*, q).

Aiyagari (1994): Formal Stationary PDE System

Production side (no aggregate shocks):

$$Y = K^{\alpha} L^{1-\alpha}, \qquad r = \alpha K^{\alpha-1} L^{1-\alpha} - \delta, \qquad w = (1-\alpha) K^{\alpha} L^{-\alpha}.$$

Household problem:

$$\max_{\{c_t\}} \mathbb{E}[] \left[\int_0^{\infty} e^{-\rho t} u(c_t)\, \mathrm{d}t \right], \qquad da_t = (wz_t + ra_t - c_t)\, \mathrm{d}t, \qquad a_t \geq \underline{a}, \quad z_t \in \{z_1, z_2\}.$$

HJB:

$$\rho V(a,z) = \max_c \left\{ u(c) + V_a(a,z)\,[wz + ra - c] + \lambda(z)\big(V(a,\hat{z}) - V(a,z)\big) \right\},$$

$$u'(c^*) = V_a, \qquad s(a,z) := wz + ra - c^*(a,z).$$

KFE on $(\underline{a}, \infty)$:

$$0 = -\partial_a\big[s(a,z) g^*(a,z)\big] - \lambda(z) g^*(a,z) + \lambda(\hat{z}) g^*(a,\hat{z}).$$

Market clearing & normalization:

$$K = \sum_z \int_{\underline{a}}^{\infty} a\, g^*(a,z)\, \mathrm{d}a, \qquad L = \sum_z \int_{\underline{a}}^{\infty} z\, g^*(a,z)\, \mathrm{d}a, \qquad \sum_z \int_{\underline{a}}^{\infty} g^*(a,z)\, \mathrm{d}a = 1.$$

Unknowns: (V, g^*, r, w, K) (or (V, g^*, r) with $w = w(r)$).

Huggett vs. Aiyagari: Key Differences

	Huggett (1993)	**Aiyagari (1994)**
Economy	Endowment	Production ($Y = K^\alpha L^{1-\alpha}$)
Asset	Bond b	Capital claim a
Supply	Zero net ($\int b = 0$)	Positive ($\int a = K > 0$)
Price	Bond return q	Interest rate r, wage w
Clearing	$\int b\, g\, \mathrm{d}b = 0$	$\int a\, g\, \mathrm{d}a = K(r)$

How the distributions differ:

Huggett:

- Wealth centered around 0 (zero net supply)
- Large mass at borrowing constraint
- Symmetric-ish with constrained left tail

Aiyagari:

- Wealth shifted right ($K > 0$)
- Long right tail (wealthy agents)
- Smaller mass at constraint

Wealth Distributions: Huggett vs. Aiyagari

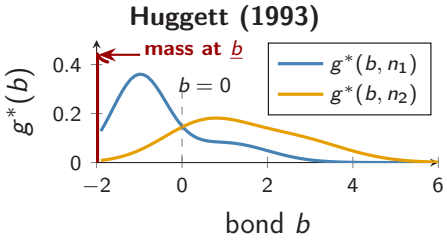

Net supply $= 0$: mean wealth ≈ 0.

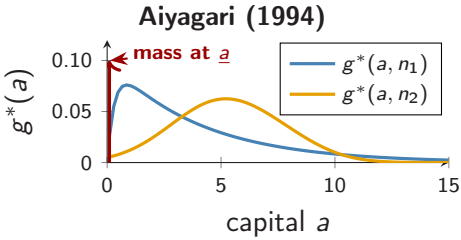

$K > 0$: mean wealth is positive with a right tail.

Why This Matters: HANK and Modern Macro

The Aiyagari/Krusell–Smith framework is the backbone of HANK:

Heterogeneous Agent New Keynesian:

- Kaplan, Moll & Violante (2018, *AER*)
- Replace representative agent with Aiyagari economy
- Add nominal rigidities
- Result: monetary policy affects consumption heterogeneously

Applications:

- Fiscal stimulus (who spends?)
- Inequality dynamics
- Wealth taxation
- COVID-19 policy responses

Key insight from HANK:

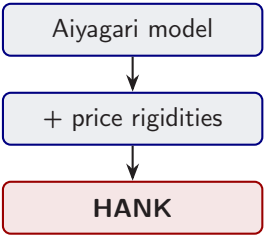

"These building blocks are essential for modern macro."

Part V

The Master Equation

From Coupled System to Single PDE

Recursive Representation of Equilibrium

States: individual $x = (a, n)$, aggregate (z, g); value function $V(x, z, g)$.
Given a belief drift $\hat{\mu}^g$, the household HJBE is

$$0 = \max_c \Big\{ -\rho V + u(c) + V_a \left[w(z,g)n + r(z,g)a - c \right] + \lambda(n)\big(V(a, \hat{n}, z, g) - V(a, n, z, g)\big)$$
$$+ V_z\, \mu^z(z) + \tfrac{1}{2}(\sigma^z)^2 V_{zz} + \sum_{j \in \{1,2\}} \int_{\underline{a}}^{\infty} \frac{\delta V}{\delta g_j}(a, n, z, g)(y)\, \hat{\mu}_j^g(y, z, g)\, \mathrm{d}y \Big\}.$$

For $c^*(a, n, z, g; \hat{\mu}^g)$, the density drift is

$$\mu_j^g(a, z, g) = -\partial_a\big[(w(z,g)n_j + r(z,g)a - c_j^*(a, z, g))\, g_j(a)\big] - \lambda_j g_j(a) + \lambda_{\hat{j}} g_{\hat{j}}(a).$$

Equilibrium consistency: $\hat{\mu}^g = \mu^g$.

The Master Equation

Key idea: Substitute KFE, market clearing, and belief consistency into the HJBE. Result: one PDE characterizing the equilibrium value function:

Master Equation

$$
\begin{aligned}
0 = -\rho V(a,n,z,g) + u(c^*) + V_a\left[w(z,g)n + r(z,g)a - c^*\right] \\
+ \lambda(n)(V(a,\hat{n},z,g) - V(a,n,z,g)) + V_z\,\mu^z(z) + \tfrac{1}{2}(\sigma^z)^2 V_{zz} \\
+ \sum_{j\in\{1,2\}} \int_{\underline{a}}^{\infty} \frac{\delta V}{\delta g_j}(a,n,z,g)(y)\,\mu_j^g(y,z,g)\,\mathrm{d}y \;=:\; \mathcal{L}V.
\end{aligned}
$$

where c^* satisfies $u'(c^*) = \partial_a V(a,n,z,g)$.

"One equation to rule them all"

HJB + KFE + market clearing = master equation

The Envelope Condition: Working with $W := \partial_a V$

Practical insight (Gu, Lauriere, Merkel & Payne, 2024): It is easier to work with $W(a, l, z, g) := \partial_a V(a, l, z, g)$ rather than V itself.
Why? Because W directly gives the consumption policy:
$u'(c^*(a, n, z, g)) = W(a, n, z, g) \;\Leftrightarrow\; c^* = W^{-1/\gamma}$
The **master equation for** W (using the envelope theorem):

$$\boxed{0 = (r(z, g) - \rho)\, W + \underbrace{\mathcal{L}_x W}_{\text{individual}} + \underbrace{\mathcal{L}_z W}_{\text{aggregate TFP}} + \underbrace{\mathcal{L}_g W}_{\text{distribution}} =: \mathcal{L} W}$$

where:

- $\mathcal{L}_x W$: individual state dynamics (savings, income switching)
- $\mathcal{L}_z W$: aggregate TFP dynamics ($\partial_z W \cdot \mu^z + \frac{1}{2}(\sigma^z)^2 \partial_{zz} W$)
- $\mathcal{L}_g W$: distribution dynamics (functional derivative w.r.t. g)

Why the Master Equation?

Advantages:

- ✓ Single PDE instead of coupled HJB + KFE + MC
- ✓ No separate market clearing step
- ✓ Amenable to neural network approximation: parameterize V or W by a deep net
- ✓ Auto-differentiation computes all derivatives (including functional derivative w.r.t. g)

Challenges:

- ✗ Infinite-dimensional: functional derivative $\frac{\delta V}{\delta g_j}$
- ✗ Need finite-dimensional approximation of g
- ✗ Training can be unstable ("bad approximate solutions")
- ✗ No convergence guarantees

Deck 2 (after coffee): Three approaches to approximate g in finite dimensions: (1) Finite population, (2) Discrete state space, (3) Projection onto basis functions. Then: solve the master equation with EMINNs (neural networks + auto-diff).

Summary and Key Takeaways

Stochastic Calculus

Brownian motion, Itô processes, Itô's lemma, Poisson jumps — the language of continuous-time models.

KFE / Fokker–Planck

Tracks density evolution of a stochastic process. In economics: the wealth distribution.

HJB + Equilibrium

Optimal control $\Rightarrow$ HJB. Coupled with KFE via savings policy. Master equation unifies all.

The road ahead:

1. **Deck 2, Part I**: Solve coupled HJB+KFE by finite differences (benchmark)
2. **Deck 2, Part II**: Solve with PINNs (mesh-free, auto-diff)
3. **Deck 2, Part III**: Solve the master equation with EMINNs (deep learning)

Coffee break!
Back at 10:45 for numerical methods.

From Theory to Code (numerics notebook)

After this theory deck, implement directly in:

- `lectures/day6/code/08_Aiyagari_Continuous_Time_FD_and_PINN_PyTorch.ipynb`

Mapping from this deck to notebook sections:

- HJB + KFE PDE system (Huggett/Aiyagari) $\rightarrow$ FD block (Section 1)
- Stationary distribution and boundary conditions $\rightarrow$ FD diagnostics + boundary plots
- PINN loss (HJB residual, KFE residual, BC terms) $\rightarrow$ PINN block (Section 2)
- Equilibrium moments and market clearing $\rightarrow$ FD/PINN comparison metrics

References

Foundational:

- Aiyagari, S.R. (1994). Uninsured Idiosyncratic Risk and Aggregate Saving. *QJE* 109(3), 659–684.
- Bewley, T. (1980). The Optimum Quantity of Money. In *Models of Monetary Economies*. Fed. Reserve Bank Minneapolis.
- Huggett, M. (1993). The Risk-Free Rate in HA Economies. *JEDC* 17(5–6), 953–969.
- Krusell, P. & Smith, A.A. (1998). Income and Wealth Heterogeneity. *JPE* 106(5), 867–896.

Continuous-time formulations:

- Achdou, Y., Han, J., Lasry, J.-M., Lions, P.-L. & Moll, B. (2022). Income and Wealth Distribution in Macroeconomics: A CT Approach. *REStud* 89(1), 45–86.
- Gu, Z., Lauriere, M., Merkel, S. & Payne, J. (2024). Global Solutions to Master Equations for CT HA Models. arXiv:2406.13726.

Mean field games:

- Lasry, J.-M. & Lions, P.-L. (2007). Mean Field Games. *Jpn. J. Math.* 2(1), 229–260.

HANK and applications:

- Kaplan, G., Moll, B. & Violante, G.L. (2018). Monetary Policy According to HANK. *AER* 108(3), 697–743.
- Gabaix, X., Lasry, J.-M., Lions, P.-L. & Moll, B. (2016). The Dynamics of Inequality. *Econometrica* 84(6), 2071–2111.
- Ahn, S., Kaplan, G., Moll, B., Winberry, T. & Wolf, C. (2018). When Inequality Matters for Macro. *NBER Macro Annual* 32(1), 1–75.

Deep learning for het. agent models:

- Fernandez-Villaverde, J., Hurtado, S. & Nuno, G. (2023). Financial Frictions and the Wealth Distribution. *Econometrica* 91(3), 869–901.
- Nuno, G. & Moll, B. (2018). Social Optima in Economies with HA. *RED* 28, 150–180.

Questions?

Simon Scheidegger

HEC, University of Lausanne

https://sischei.github.io

Materials: https://github.com/sischei/Deep_Learning_for_Solving_And_Estimating_Dynamic_Economic_Models

Covered today (Deck 1):

I	Stochastic Calculus
II	Kolmogorov Forward Eq.
III	Hamilton–Jacobi–Bellman
IV	Competitive Equilibrium
V	The Master Equation

Next up (Deck 2):

Numerical Methods

FD $\rightarrow$ PINNs $\rightarrow$ EMINNs

10:45–12:00